DEV CON VIII INDIA

MUMBAI, INDIA
3—6 Nov, 2026

What Actually Happens When You Send an LLM Request to Kubernetes?

Atharva Mhaske



Backend Engineer at Cloud9 · Open Source Member at OpenEverest, leading provider-kserve

CLS – LIGHTNING TALK OTS

foss asia ots OpenTechSummit India · Devcon 8

The problem

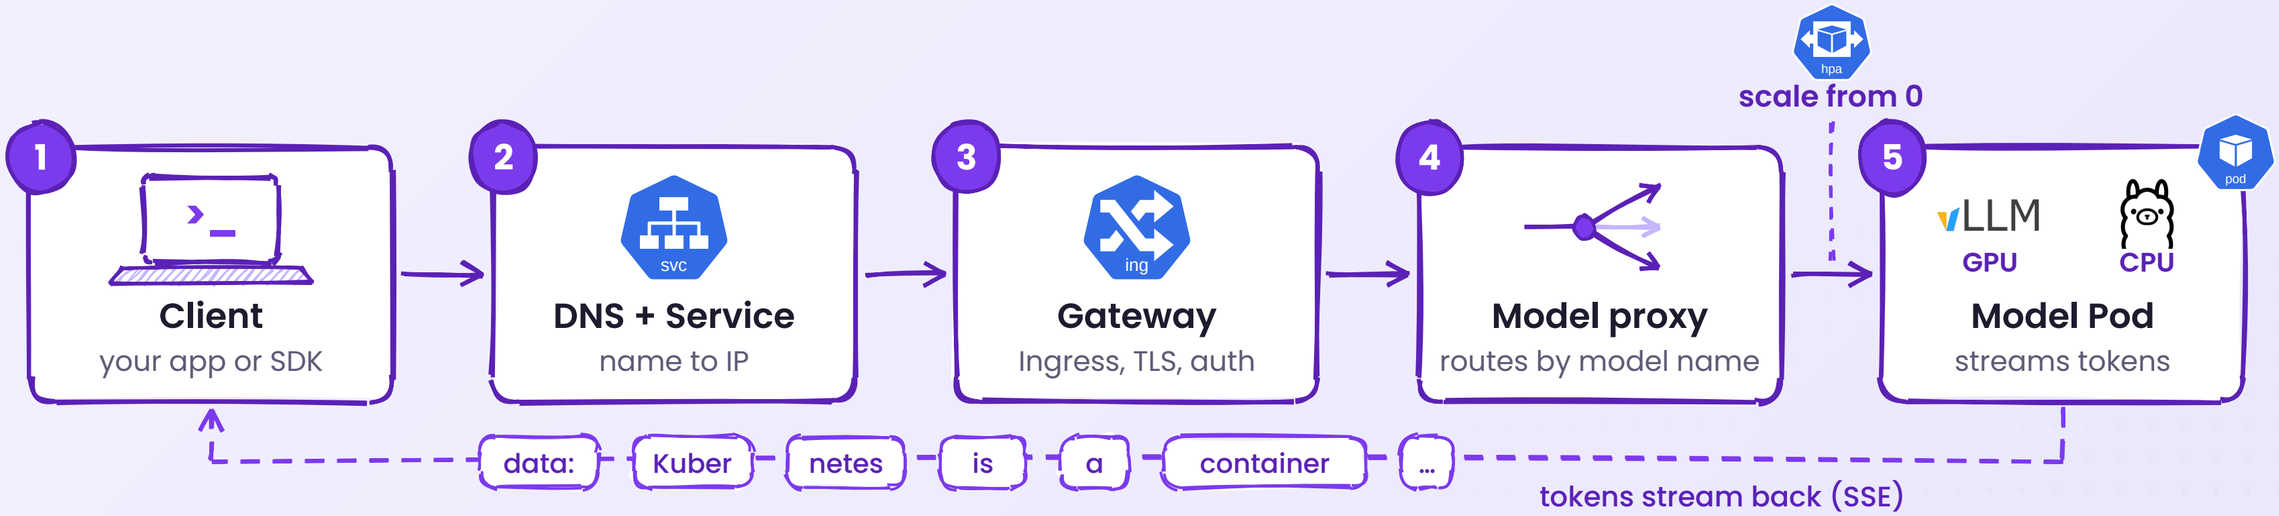
You called /v1/chat/completions. Then what?

Platform teams can't see the path, so LLM calls fail in ways nobody can explain.



The path

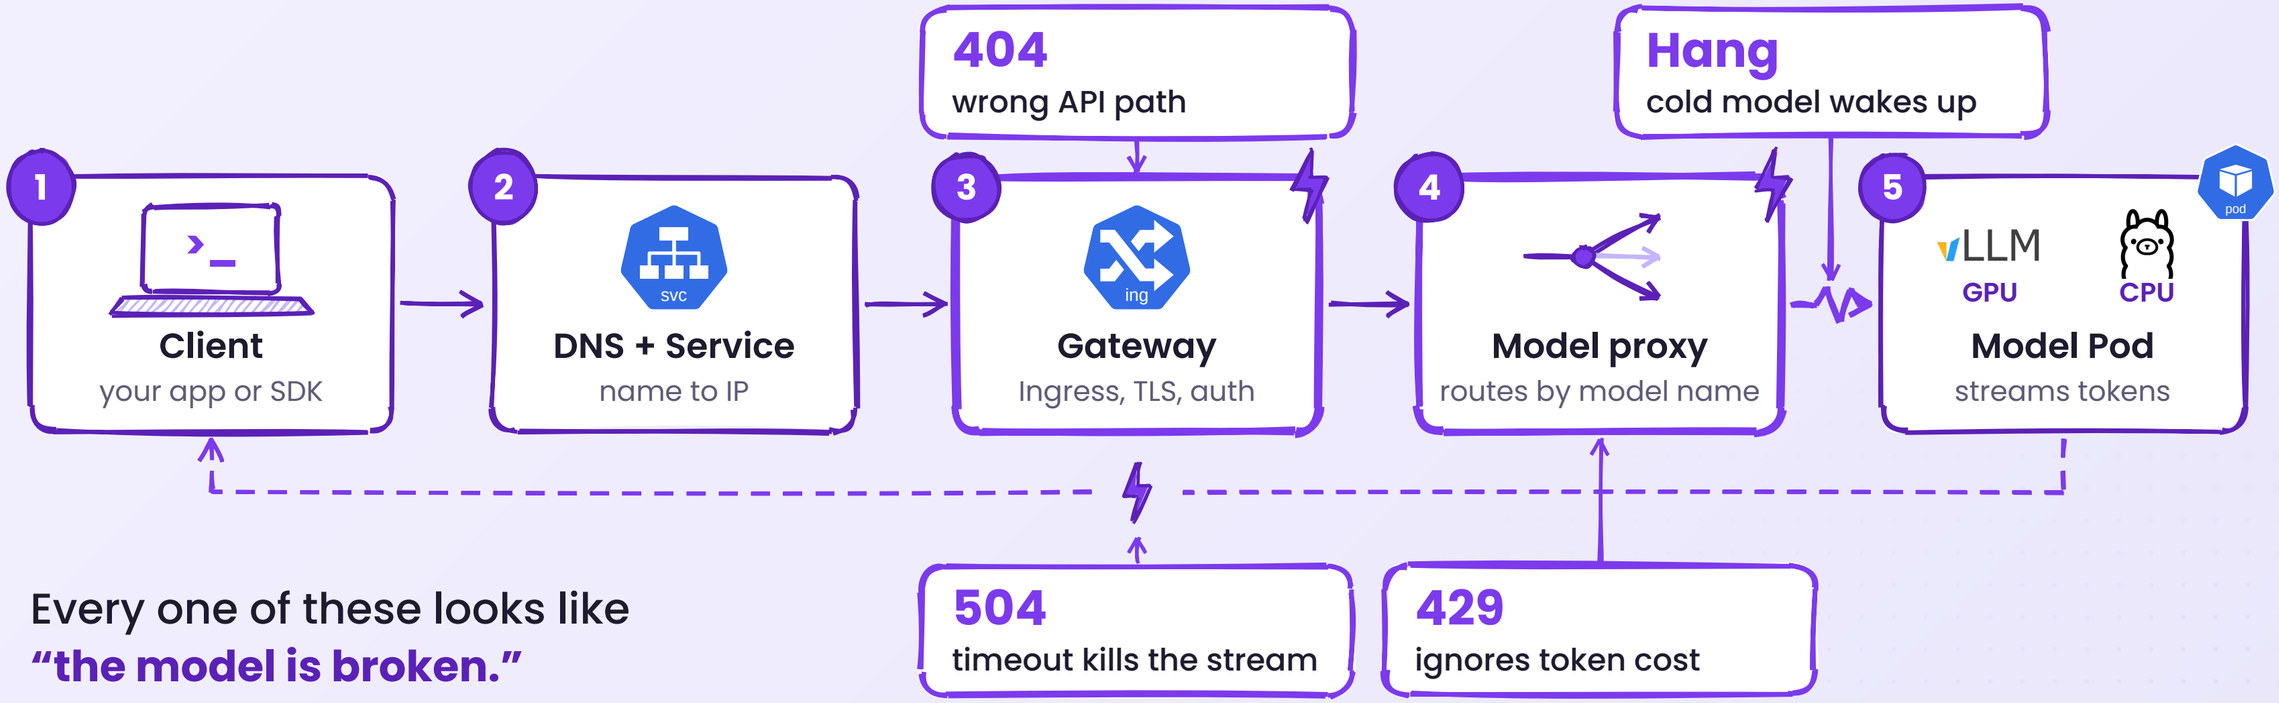
Follow one request, hop by hop



It’s never one Pod.

Where it breaks

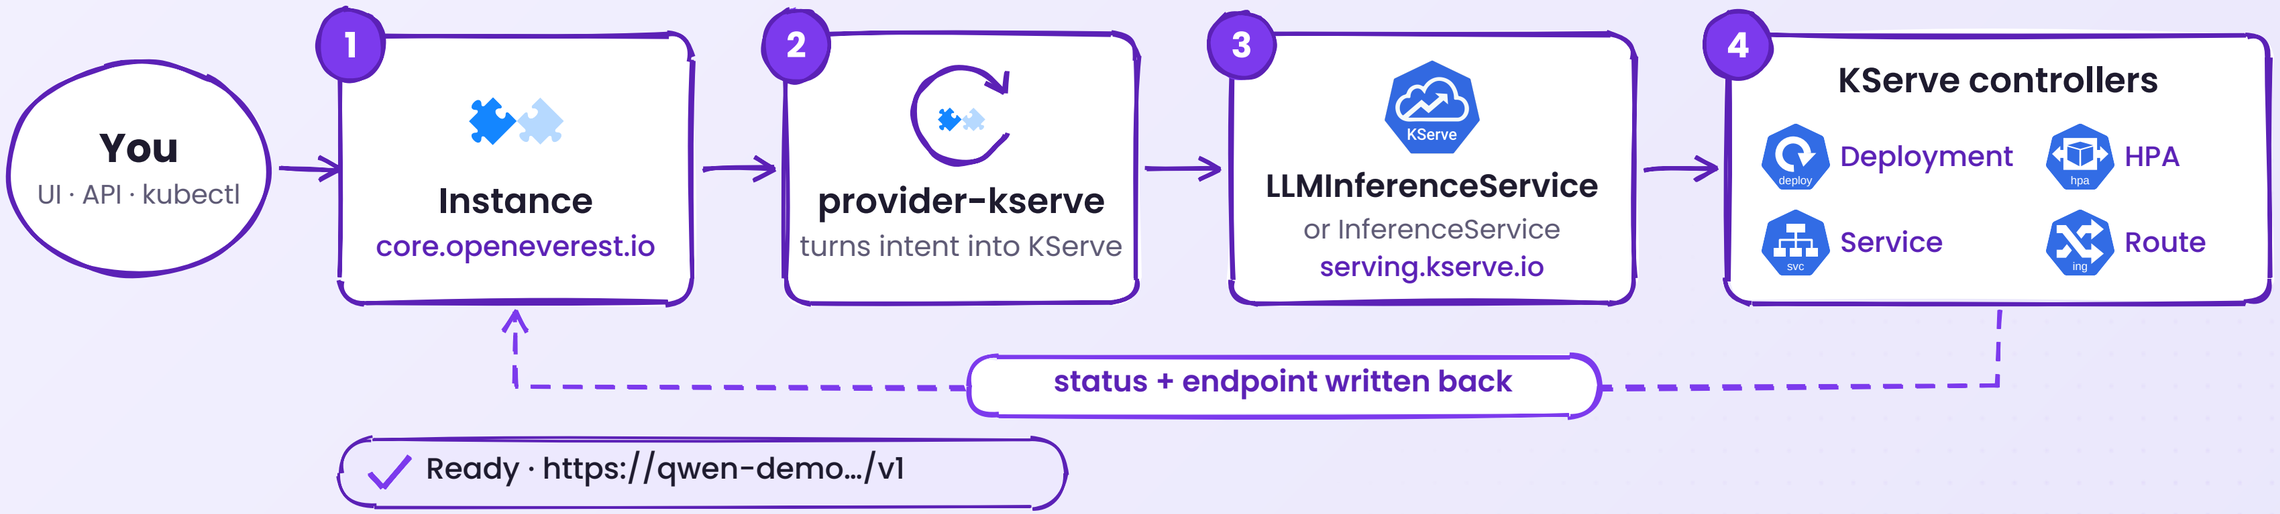
Each hop fails in its own way



Every one of these looks like
“the model is broken.”

How OpenEverest solved it

One object in, a working endpoint out



Declare intent. Get an endpoint back.

Same hops, any stack

Different tools, the same four hops

	KServe	KubeAI	Hand-rolled
Gateway	Envoy AI Gateway	Ingress to KubeAI	ingress-nginx
Model routing	InferencePool + endpoint picker	proxy reads the *model* field	one path per model
Scale from zero	Knative or KEDA	built in, to zero	HPA, min 1 replica
Model Pod	vLLM runtime	vLLM or Ollama	vLLM Deployment

Map the path, not the brand.

DEV CON VIII INDIA

MUMBAI, INDIA · 5 NOV 2026

OTS Main Stage · 16:05 IST

Thank you



Atharva Mhaske

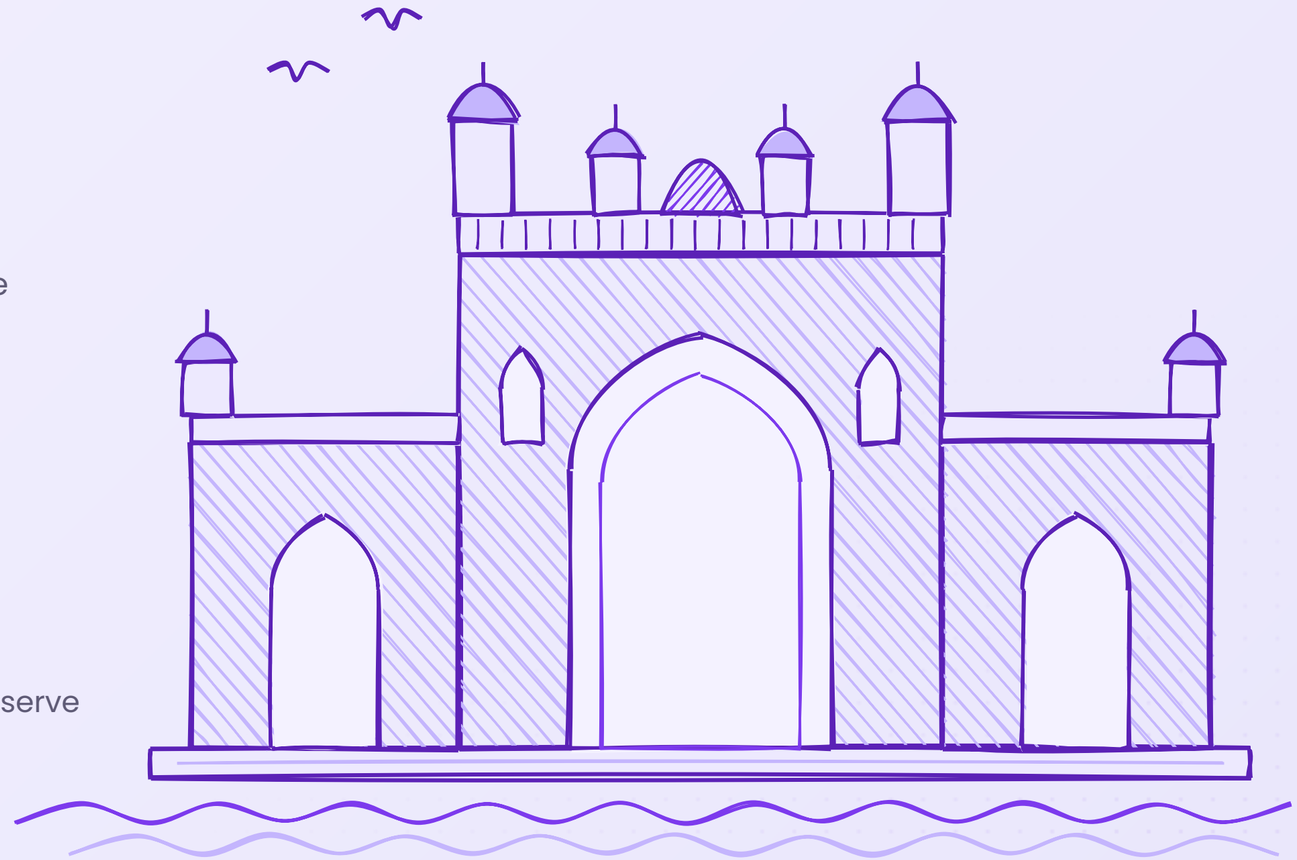
Backend Engineer, Cloud9 · OpenEverest, provider-kserve

Scan for my speaker page

eventyay.com/fossasia/ots2026/speakers/ZNTNQ3

github.com/atharvamhaske

Try it: github.com/openeverest/provider-kserve



Gateway of India, Mumbai

FOSSASIA OTS | Open Everest