DEV CON VIII INDIA

MUMBAI, INDIA
5 Nov, 2026
OTS Main Stage · 16:15 IST

What Actually Happens When You Send an LLM Request to Kubernetes?



Atharva Mhaske
Engineer at Cld9.ai

CLS – LIGHTNING TALK | OTS

foss asia | Open Everest

Who am I

Atharva Mhaske

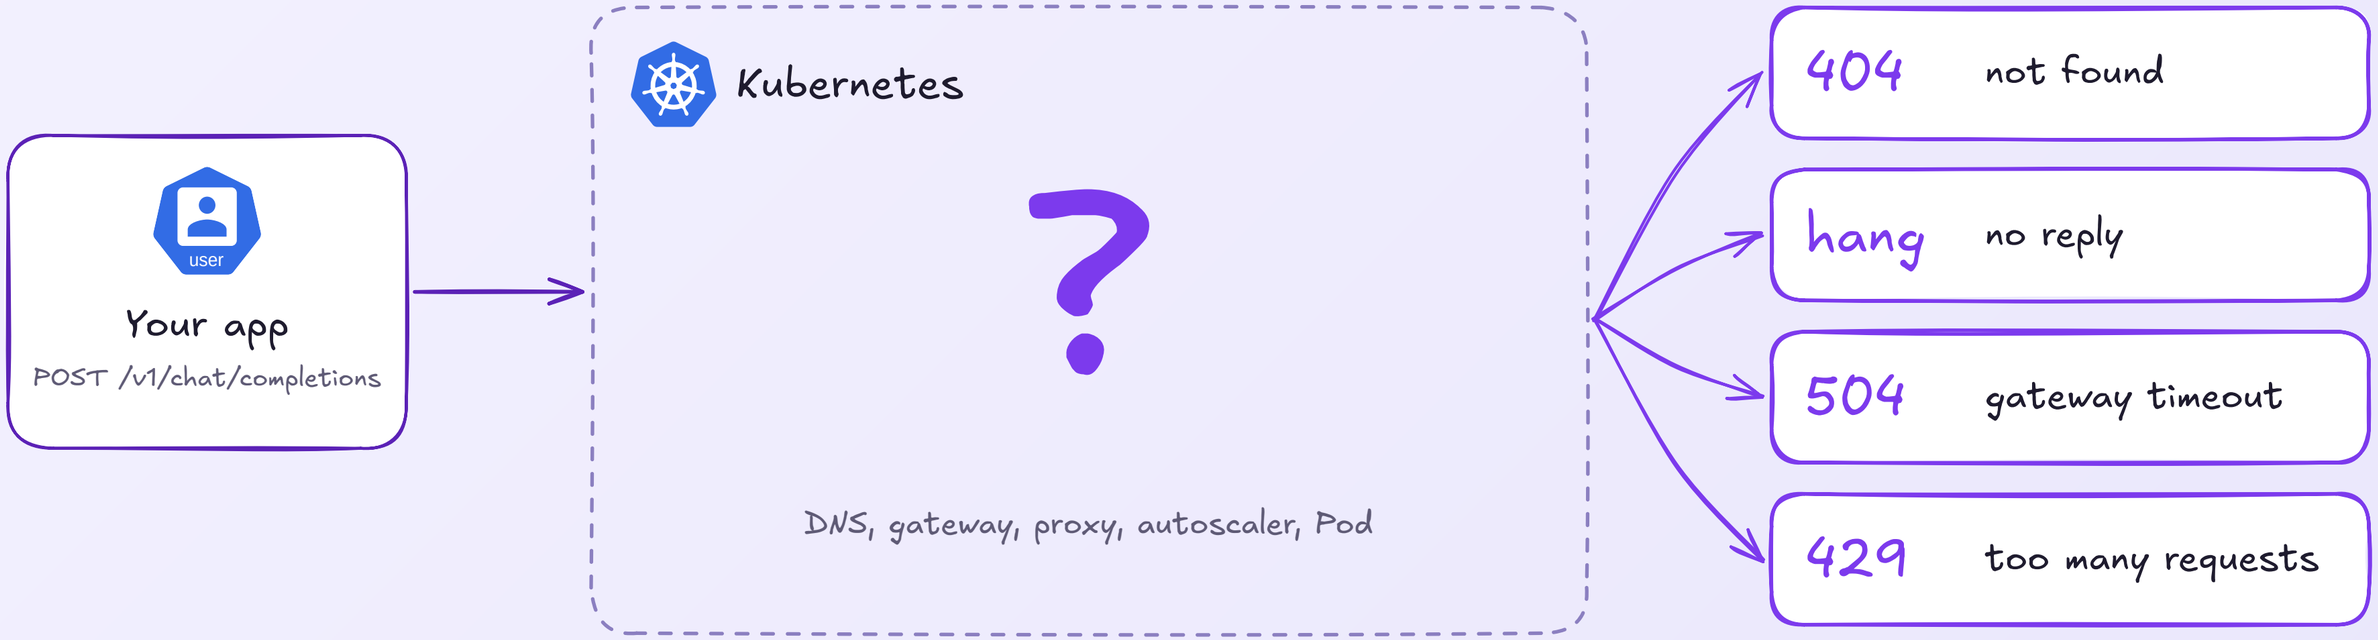
- AI and Infra Engineer at Cld9.ai
- Former Google Summer of Code 2026 Mentee
- Member @OpenTelemetry project
- I love FOSS and contribute to multiple OSS projects under CNCF



THE PROBLEM

Where did this request go?

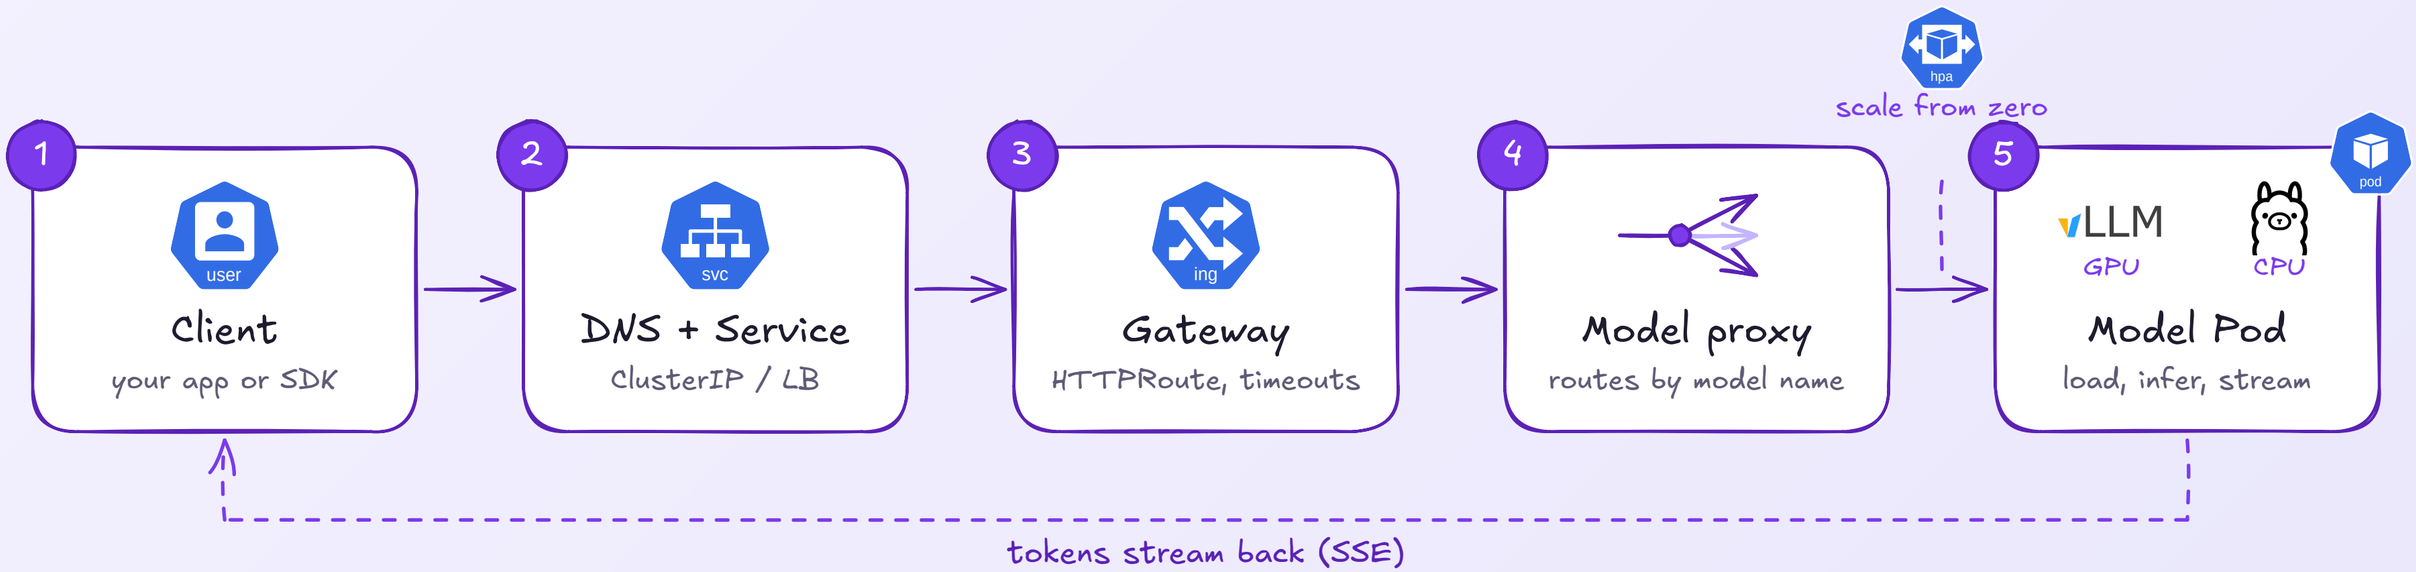
Platform teams ship LLM features fast. In production, one chat completion is rarely one Pod.



Teams debug the symptom because they never mapped the path.

THE PATH

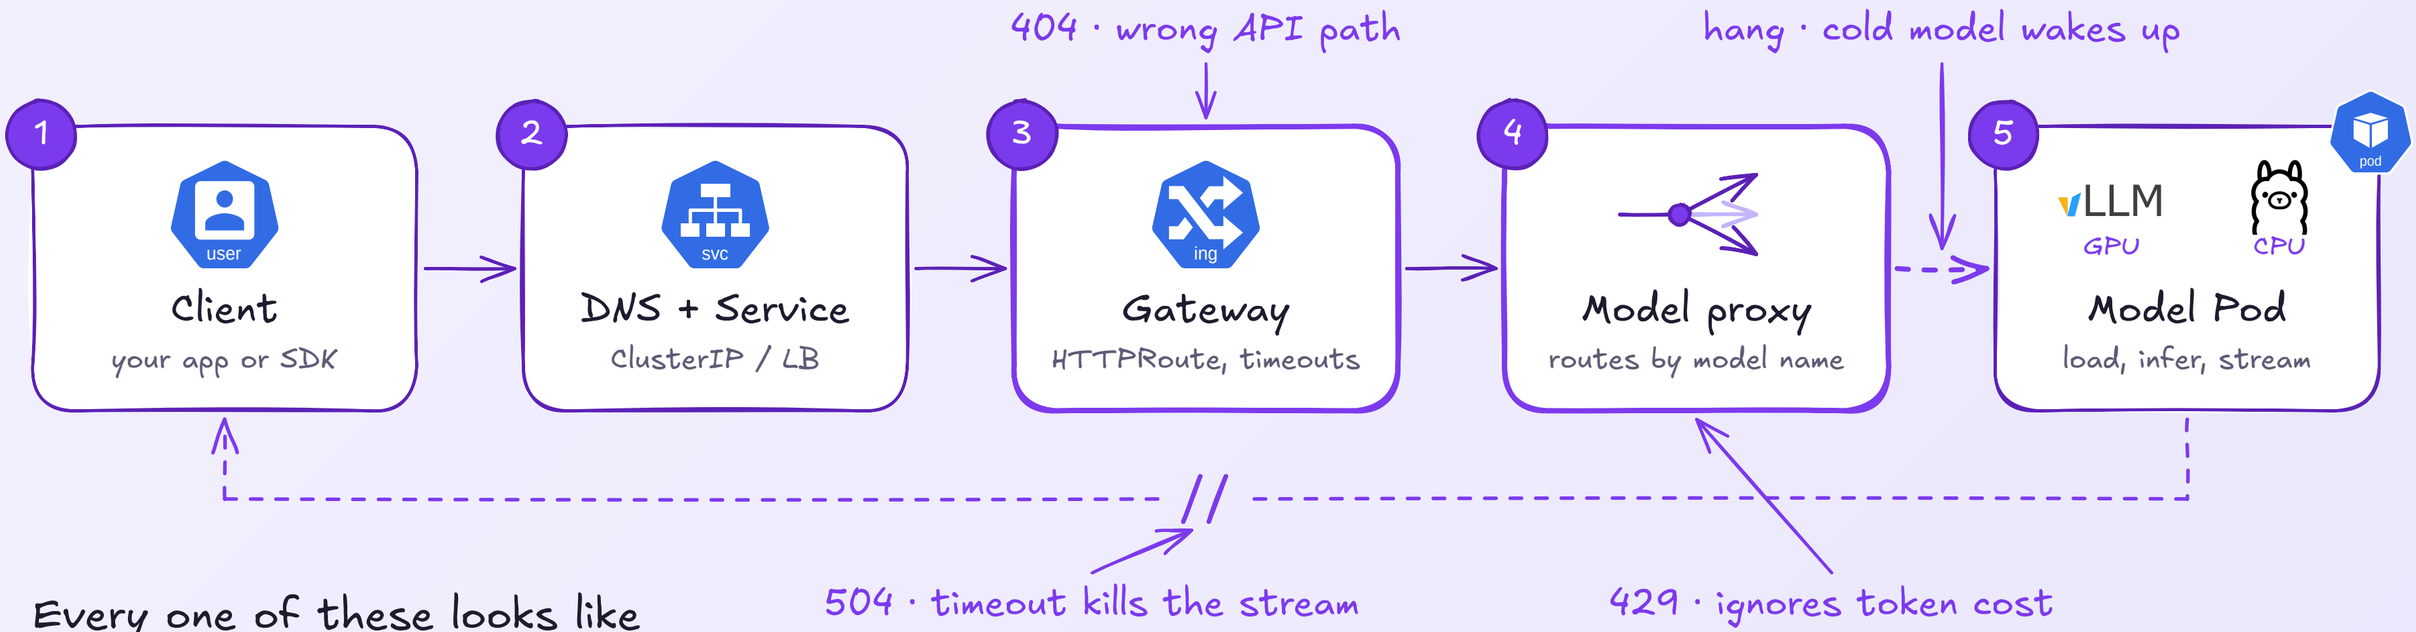
Follow one request, hop by hop



It's never one Pod.

WHERE IT BREAKS

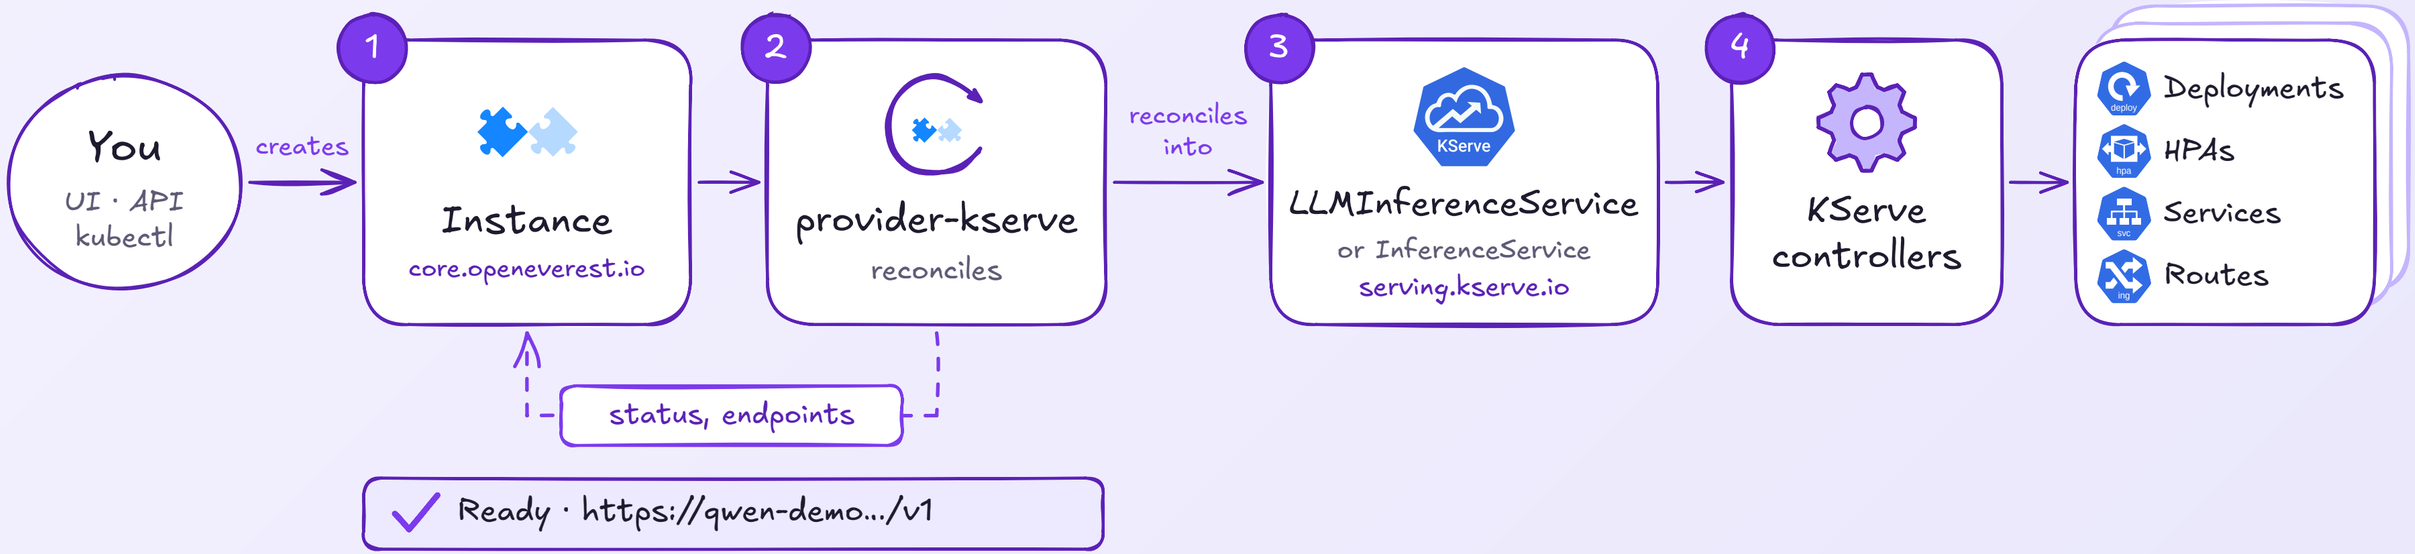
Each hop fails in its own way



Every one of these looks like "the model is broken."

HOW OPENEVEREST SOLVED IT

provider-kserve turns an Instance into an endpoint



Declare intent. Get an endpoint back.

SAME HOPS, ANY STACK

Three stacks, the same four hops

	KServe	KubeAI	Hand-rolled
Gateway	Envoy AI Gateway	Ingress to KubeAI	ingress-nginx
Model routing	InferencePool + endpoint picker	proxy reads the model field	one path per model
Scale from zero	Knative or KEDA	built in, to zero	HPA, min 1 replica
Model Pod	vLLM runtime	vLLM or Ollama	vLLM Deployment

Map the path, not the brand.

Thank you

Let's connect after the talk.

Scan to connect on LinkedIn

linkedin.com/in/atharavxdevs



FOSS asia | Open Everest